

Machine Learning Basics

Prof. Artur Jordão

The Many Faces of Artificial Intelligence

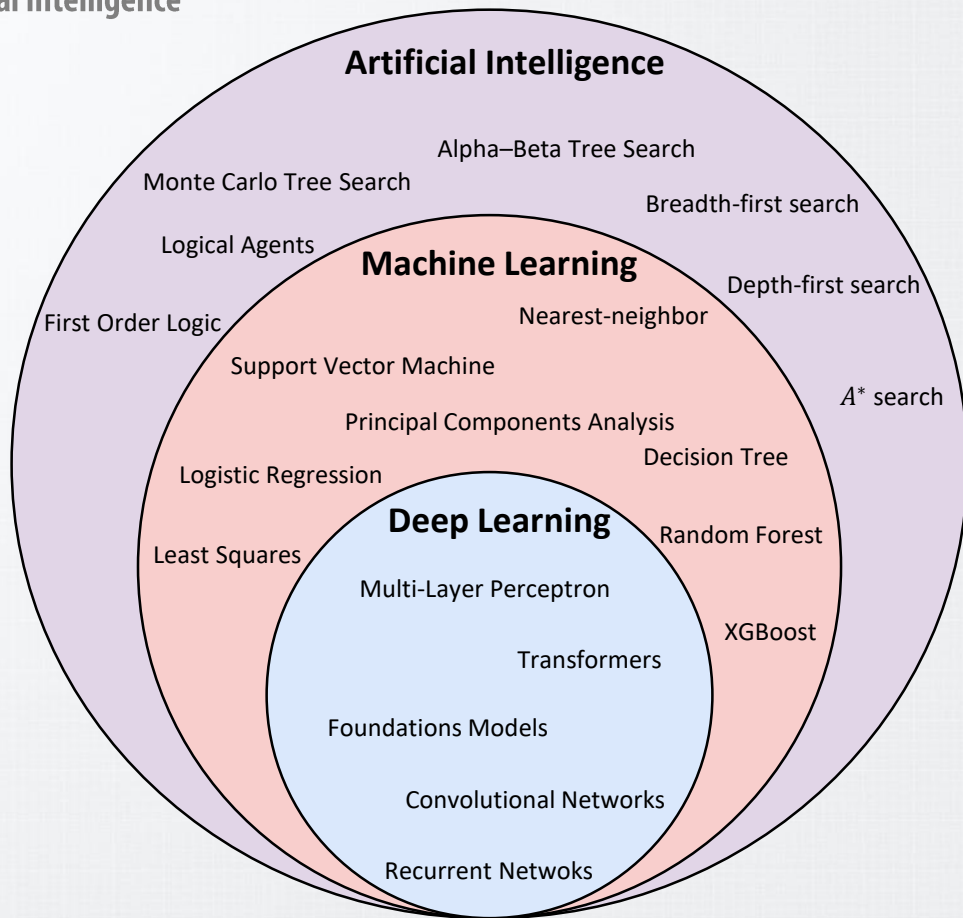
Definitions

The Many Faces of Artificial Intelligence

- Artificial Intelligence
 - Any technique that enables computers to mimic human behavior
- Machine Learning
 - Ability to learn without explicitly being programmed
- Deep Learning
 - Extract patterns from data using neural networks

Machine Learning and Deep Learning

The Many Faces of Artificial Intelligence

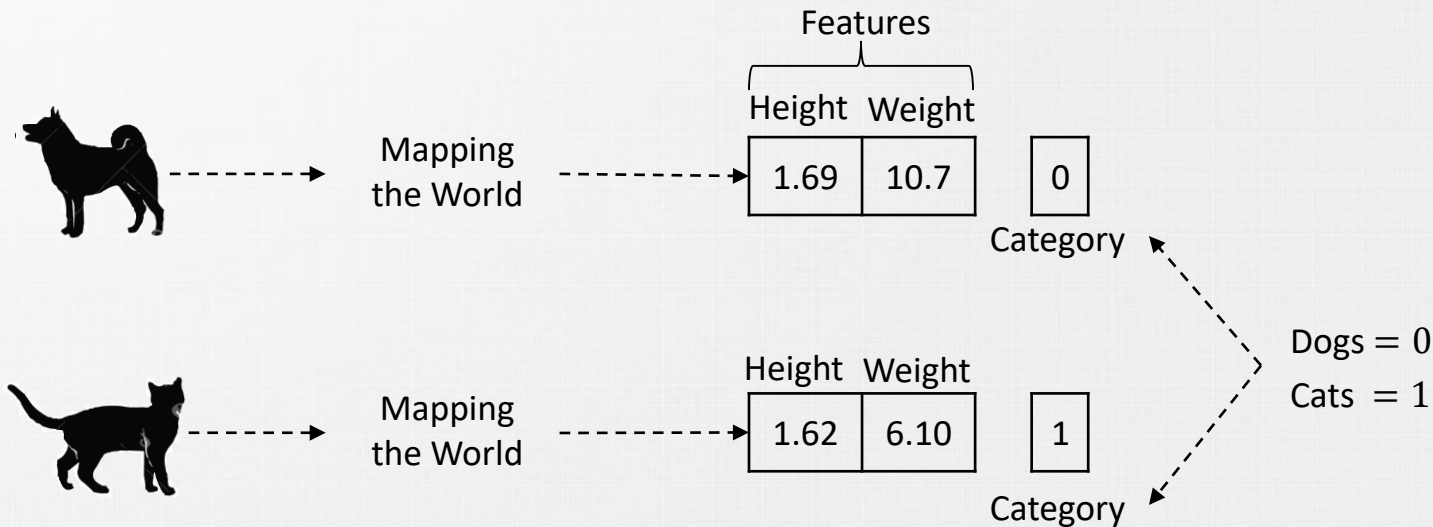


Independent (X) and Dependent (Y) Variables

Preliminaries

Independent and Dependent Variables

- Mapping the state of the world (features – m)
 - How can we represent the states (i.e., objects) of the world?
 - Numerical representation (features engineering)



Preliminaries

Independent and Dependent Variables

- Assume that we represent (map) n dogs and cats using m features
- Let $X \in \mathbb{R}^{n \times m}$
 - Independent variables
 - Data samples – each dog/cat composes a row in X
- Let $Y \in \mathbb{R}^{n \times c}$
 - c stands for the number of categories (i.e., 2)
 - Dependent variable
 - Labels/Classes

	X		Y
n	1.69	10.7	0
	1.85	14.5	0
	2.67	20.3	0

	1.62	6.1	1
	1.11	4.0	1
	m		

Preliminaries

Independent and Dependent Variables


- Let $x_i \in \mathbb{R}^{1 \times m}$ be the i th sample (example) of X
 - We can express x_i in terms of its features $x_i = x_i^1, x_i^2 \dots x_i^m$
- Let $y_i \in \mathbb{R}^{1 \times 1}$ be the i th label of Y

X		Y
1.69	10.7	0
1.85	14.5	0
x_i 2.67	20.3	0 y_i
....
1.62	6.1	1
1.11	4.0	1

Multiclass Problems

Independent and Dependent Variables

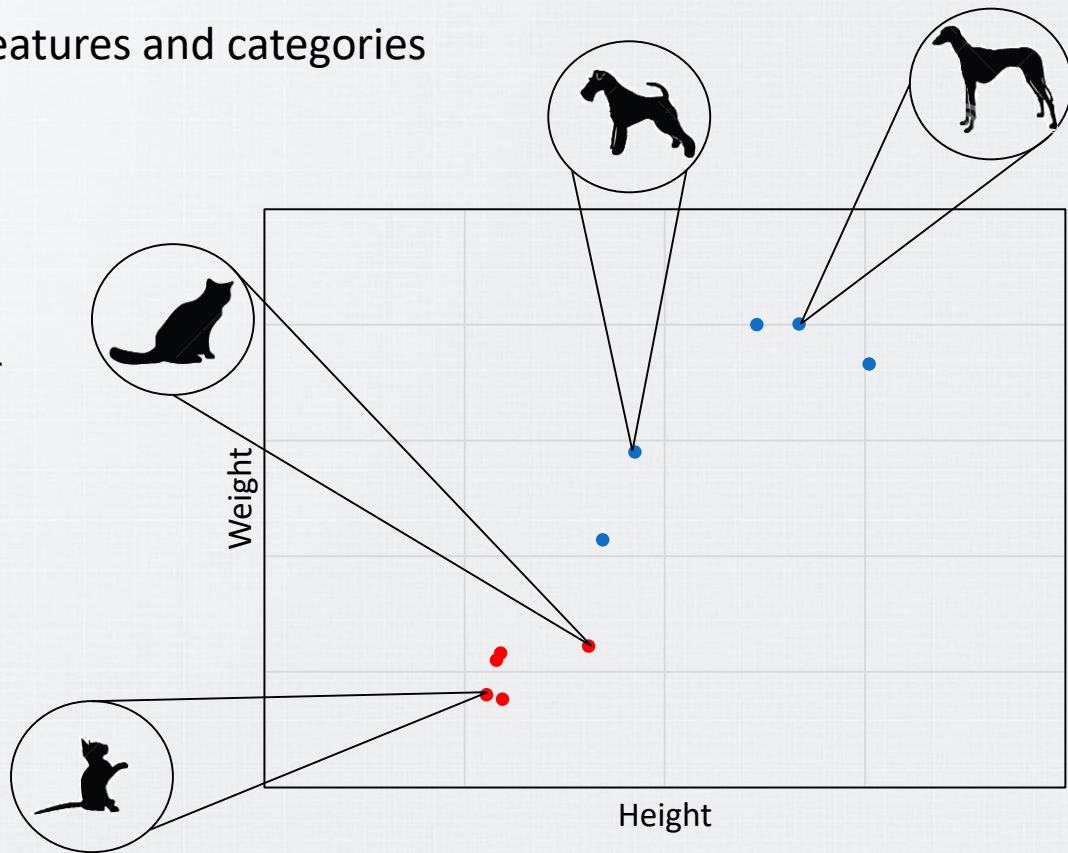
- Problems involving more than two classes
 - In this case, we cannot use 0 and 1
 - For example, dogs, cats and bears
- One-hot encoding
 - Transform the c classes into a zero-one vector
 - The cth entry equal to 1 and the rest 0
- $Y \in \mathbb{R}^{n \times c}$



Feature Space

Independent and Dependent Variables

- Projects the data using its features and categories
- Technical details
 - Each feature (m) is an axis
 - Each sample (n) is a point
 - Each category (y) is a color



Problem Statement

Independent and Dependent Variables

- Given X , the problem is to predict Y to new samples
 - Given the features, we want to predict the category to which samples belong
- Assume $\mathcal{F}(\cdot, \cdot)$ a model parameterized by a set of parameters/weights θ that receives x and outputs \bar{y}
 - Formally, $\bar{y} = \mathcal{F}(x, \theta)$ (or $\bar{Y} = \mathcal{F}(X, \theta)$)
- The problem is, therefore, to discover θ**

X		\bar{Y} (unseen)
1.33	8.7	?
2.78	3.5	?
2.29	7.3	?
....
1.62	9.1	?
2.15	10.0	?

Problem Statement

Independent and Dependent Variables

- Training (or learning) phase
 - Estimate θ using X and Y (seen)
- Testing phase
 - Employ θ onto X to predict \bar{Y} (unseen)
 - Unseen means new samples

X		Y (seen)
1.69	10.7	0
1.85	14.5	0
....
1.62	6.1	1
1.11	4.0	1

Training Phase

X		\bar{Y} (unseen)
1.33	8.7	?
2.78	3.5	?
....
1.62	9.1	?
2.15	10.0	?

Testing Phase

Classification vs. Regression

Independent and Dependent Variables

- Classification
 - The goal is to predict a category between c possible categories
- Regression
 - The goal is to predict a real-valued target
- For both classification and regression, the problem is to discover θ
 - $\bar{Y} = \mathcal{F}(X, \theta)$


X (features)	Task	Y (target)
Weight, Height (Animal features)	Classification	Dogs, Cats, Bears
	Regression	Lifespan
Area, location, number of bedrooms (House features)	Classification	House Quality (good, bad, medium)
	Regression	House Price

Preliminaries

Functions

Preliminaries



- A function is a relation that associates each element x to **a single** element $f(x)$
- Global minimum
- Local minimum



Convex and Non-Convex Problems

Preliminaries


- Convex Problems: **only one** global minimum (or maximum)
 - It facilitates the optimization process (i.e., discovering θ)
- Non-Convex Problems: multiple local minimum (or maximum)
 - Multiple optimal local solutions
 - Most practical problems belong to this category – deep learning problems



Loss Function

Preliminaries

- The function we want to minimize or maximize is called the objective function or criterion
 - When we are minimizing a function, we may also call it the **cost function**, **loss function**, or **error function**
- We often denote the value that minimizes (or maximizes) a function with a superscript *
 - $x^* = \operatorname{argmin} f(x)$



Loss Function

Preliminaries

- Quantify the distance between the real (Y) and predicted values of the target
 - Real values: y or Y
 - Predicted values: $\mathcal{F}(x, \theta) = \hat{y}$ or $\mathcal{F}(X, \theta) = \hat{Y}$
- Loss function properties
 - Monotonicity: The better the model gets, the lower the value of the loss function
 - Differentiability: Differentiable with respect to θ

Common Losses	Short Description
Squared error	$\frac{1}{2}(\hat{y} - y)^2$
Mean Square Error	$\frac{1}{n} \sum_{i=0}^n (\hat{y}_i - y_i)^2$
Cross Entropy	$-\sum_{i=0}^c y_i \log(\hat{y}_i)$


The Role of the Derivative

Preliminaries

- Define a function $y = f(x)$ and $f'(x)$ the derivative of f
 - x and y are both real numbers
- The derivative $f'(x)$ gives the slope of $f(x)$ at the point x
 - It specifies how to scale a small change in the input in order to obtain the corresponding change in the output: $f(x + \epsilon) \approx f(x) + \epsilon f'(x)$

$$\frac{\partial f}{\partial x} = \lim_{h \rightarrow 0} \frac{f(x+h) - f(x)}{h}$$

Quantity $\rightarrow \frac{\partial f}{\partial x}$
 w.r.t $\rightarrow \frac{\partial f}{\partial x}$
 Time interval $\rightarrow h$
 Change in the value of quantity $\rightarrow f(x+h) - f(x)$
 New value $\rightarrow f(x+h)$
 Old value $\rightarrow f(x)$
 Tends to zero (becomes very small) $\rightarrow h \rightarrow 0$



The Role of the Derivative

Preliminaries

- The **gradient** generalizes the notion of derivative to the case where the derivative is with respect to a vector
- The gradient of f is the vector containing all of the partial derivatives $\nabla_x f(x)$
 - Element i of the gradient is the partial derivative of f with respect to x_i

Linear Regression

Regression

Linear Regression

- Regression refers to a set of methods for modeling the relationship between one or more independent variables and a dependent variable
- Linear regression
 - The simplest and most popular among the standard tools for regression

Assumptions of Linear Regression

Linear Regression

- Linearity assumption
 - The relationship between the independent variables x and the dependent variable y is linear
 - y can be expressed as a **weighted sum** of the elements in x , given some noise on the observations
- Any noise is well-behaved
 - Following a Gaussian distribution

Linear Model

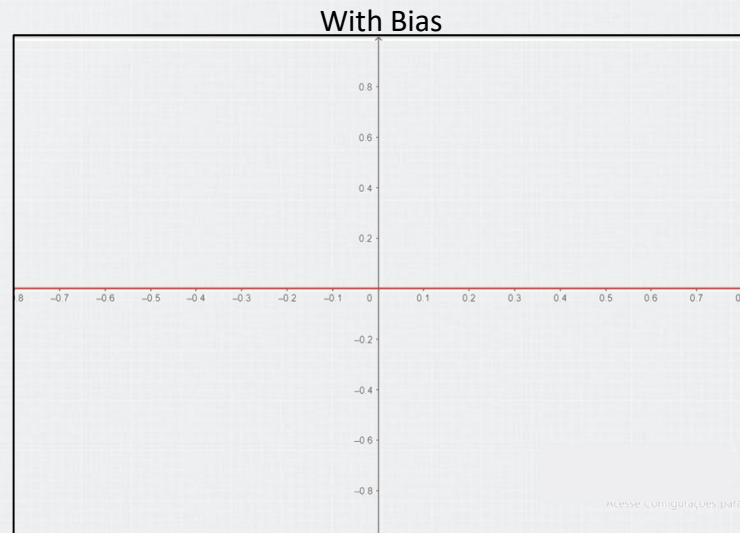

Linear Regression

- $\hat{y} = w_1 * x^1 + w_2 * x^2 + \dots + w_m * x^m + \mathbf{b}$
 - $W = w_1, w_2, \dots, w_m$
 - We can organize W as column or row matrix
- Dot product form (single sample prediction)
 - $\hat{y} = xW^T + b$
- Matrix vector product form (n samples prediction – all at once)
 - $\hat{Y} = XW^T + b$

The Bias Term

Linear Regression

- The bias term plays a role in the expressivity of the model
 - It allows the model to fit not only the data points that **pass through the origin but also those that do not**
 - It enables the model to capture and represent more complex relationships between the features and the target variable



The Bias Term

Linear Regression

- Putting the bias into the parameter matrix w_i
 - We can subsume the bias b into the parameter matrix W by appending a column to the design matrix consisting of all **ones**

X			Y (observable)
1.69	10.7	1	0
1.85	14.5	1	0
2.67	20.3	1	0
....
1.62	6.1	1	1
1.11	4.0	1	1

X			\bar{Y} (unseen)
1.33	8.7	1	?
2.78	3.5	1	?
2.29	7.3	1	?
....
1.62	9.1	1	?
2.15	10.0	1	?

Problem Definition

Linear Regression

- How can we discover W (and b) that minimizes the total loss across all training examples?

- Formally: $W^*, b^* = \operatorname{argmin} \mathcal{L}(W, b) \Rightarrow \frac{1}{n} \sum_{i=1}^n \frac{1}{2} (\overbrace{x_i W^T + b}^{\hat{y}} - y_i)^2$
- Here $\mathcal{L}(\cdot)$ means a loss function

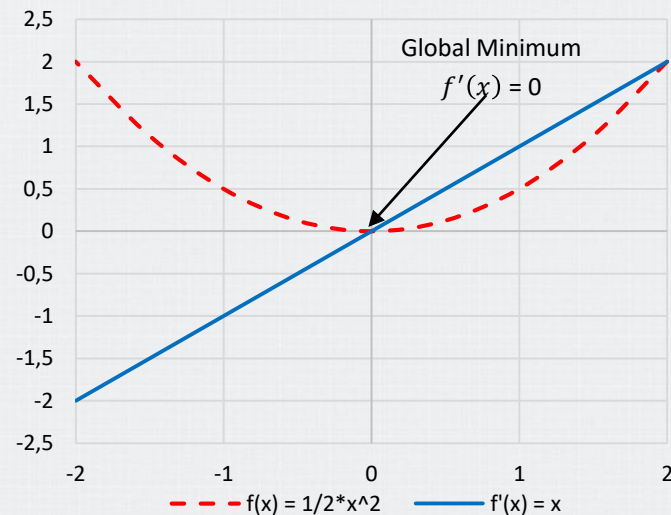
- Techniques
 - Analytic Solution
 - Gradient-Based Optimization

Analytic Solution

Linear Regression

- $\mathcal{L}(W) = ||XW^T - Y||$
 - The loss across all training samples
- $\nabla_w \mathcal{L}(W) = 2X^T(XW^T - Y) = 0$
 - We set the gradient to zero to find the points where the loss function reaches the minimum



$$W^* = \underbrace{(X^T X)^{-1} X^T}_{\text{The pseudoinverse of } X} Y$$



Analytic Solution

Linear Regression


- Problems
 - It does not work for complex (non-convex – multiple local optima) problems
- The sample size problem (a.k.a zero determinant or singularity)
 - The number of samples is smaller than the number of features ($n < m$)



Gradient Descent

Linear Regression

- The gradient is useful for minimizing a function
 - It tells us how to change x in order to make a small improvement in y
- We can reduce $f(x)$ by moving x in small steps with **opposite sign of the gradient**
- The key consists of **iteratively** reducing the error by updating the parameters (weights) in the direction that incrementally lowers the loss function
 - **Gradient Descent**



Gradient Descent

Linear Regression

- Learning rate (η)
 - Positive scalar determining the size of the step
 - The rate of learning
- Convergence
 - Run for a defined number of iterations (**epochs**)
 - Stop when the loss does not decrease (early stop)

Gradient Descent Algorithm

$\mathbf{W} \leftarrow$ Random values

While not converged do

for each $\mathbf{w}_i \in \mathbf{W}$ do

$$\mathbf{w}_i \leftarrow \mathbf{w}_i - \eta \frac{\partial}{\partial \mathbf{w}_i} \mathcal{L}(\mathbf{W})$$

Gradient Descent

Linear Regression

- The Gradient Descent takes the derivative of the loss function, which is an **average of the losses** computed on every single example ($x \in X$)
- In practice, this can be extremely slow and memory costly
 - We must pass over the entire dataset before making a single update
 - The problem is compounded if n is larger than the processor's memory size

Stochastic Gradient Descent (SGD)

Linear Regression

- Stochastic Gradient Descent randomly selects a small number (**batch size** – β) of training examples at each step t , and updates according to

$$\mathbf{W} \leftarrow \mathbf{W} - \eta \frac{1}{|\beta|} \sum_{j=\beta_t}^n \frac{\partial}{\partial \mathbf{W}} \mathcal{L}^j(\mathbf{W})$$

Stochastic Gradient Descent Algorithm

$\mathbf{W} \leftarrow$ Random values

While not converged do

 for each $\mathbf{w}_i \in \mathbf{W}$ do

$$\mathbf{w}_i \leftarrow \mathbf{w}_i - \eta \frac{1}{|\beta|} \sum_{j=\beta_t}^n \frac{\partial}{\partial \mathbf{w}_i} \mathcal{L}^j(\mathbf{W})$$

Learning and Testing Phase

Definitions

Learning and Testing Phase

- Define $D = \{(x_i, y_i)\}_{i=1}^n$ a dataset
- Let D_{train} be a subset of D such that $D_{train} \subseteq D$
 - Samples and their (seen) labels
- Let D_{test} be a subset of D such that $D_{test} \subseteq D$
 - In practice, unseen labels
- Important properties
 - $D = D_{train} \cup D_{test}$
 - $D_{train} \cap D_{test} = \emptyset$

Overview

Learning and Testing Phase

Learning Phase

(Here X and Y come from D_{train})

```
model.fit(X, Y)
```

$$(X^T X)^{-1} X^T Y$$

$$\mathbf{w}_i \leftarrow \mathbf{w}_i - \eta \frac{\partial}{\partial \mathbf{w}_i} \mathcal{L}(\mathbf{W})$$

$$\mathbf{w}_i \leftarrow \mathbf{w}_i - \eta \frac{1}{|\beta|} \sum_{j=\beta_t}^n \frac{\partial}{\partial \mathbf{w}_i} \mathcal{L}^j(\mathbf{W})$$

Testing Phase

(Here X comes from D_{test})

```
y_pred = model.predict(X)
```

$$\hat{Y} = XW^T + b$$

Quality of the Learning Trajectory

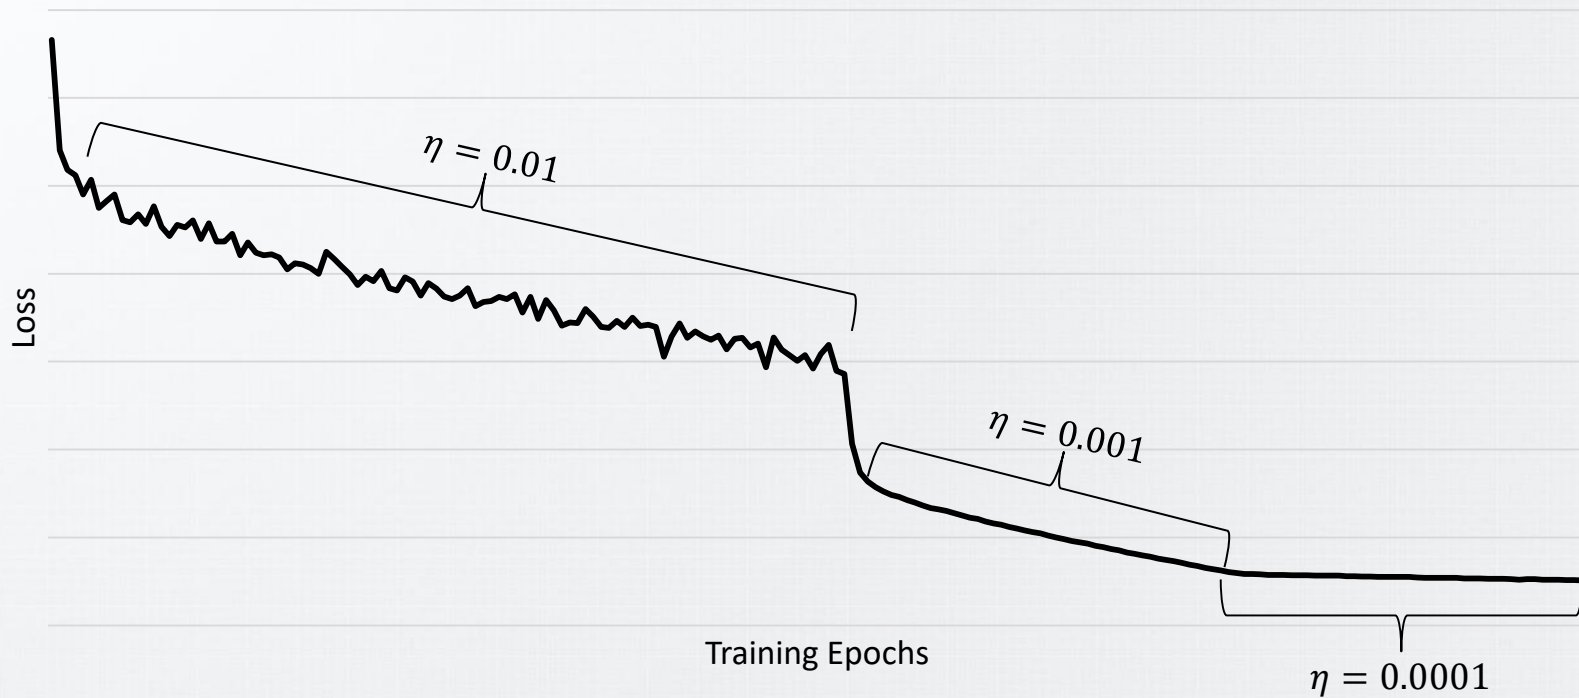
Learning and Testing Phase

- How to measure the quality of the training trajectory?
 - Dynamics of training
- Loss curve
- Loss landscape (Li et al., 2018)

Loss Curve

Learning and Testing Phase

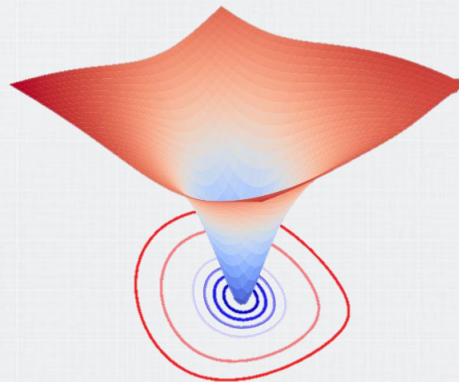



- Early stop



Loss Landscape

Learning and Testing Phase

- <http://www.telesens.co/loss-landscape-viz/viewer.html>



Testing Phase

Learning and Testing Phase

- Once the training is done, how can we measure the predictive ability of the model?
- Predictive ability (quantitative) metrics
 - Accuracy
 - Confusion Matrix
 - Loss
 - Pearson Correlation
 - Pair-wise
- The metric depends on the application
 - Some benchmarks have their own metrics
 - For example, accuracy on CIFAR-10 and Top-5 accuracy/error on ImageNet

Overfitting, Underfitting, Generalization and No Free Lunch

Generalization


Overfitting, Underfitting, Generalization and No Free Lunch

- Generalization
 - The quality of the model in predicting new (unseen) data
- Overfitting
 - The model performs well on training but poorly on testing/validation
 - Complex models tend to (we can avoid/handle this) overfit the data
- Underfitting
 - The model fails to find a pattern in the data

The Bias-Variance Tradeoff

Overfitting, Underfitting, Generalization and No Free Lunch

- Bias
 - Low bias: the model predicts well the samples of the training data
 - High bias: the model makes many mistakes in the training data
- Variance
 - Error of the model due to its sensitivity to small fluctuations in the training set
- Bias–variance tradeoff
 - Low-bias hypotheses that fit the training data well
 - Low-variance hypotheses that **may** generalize better
- U-shaped curve



No Free Lunch Theorem

Overfitting, Underfitting, Generalization and No Free Lunch

- The *no free lunch theorem* states that every learning algorithm is as good as any other when averaged over all possible problems (Wolpert, 1996)
 - **There is no universal learning algorithm able to solve all tasks precisely**
- Under a uniform distribution over problems (search/learning problems), all algorithms perform equally
 - A particular model or algorithm is **better** than average on some problems, it must be **worse** than average on others

Normalization

Z-Score

Normalization

- Suppose we are mapping the world using the following features
 - $x^1 \in [0, 100], x^2 \in [0, 1], \dots, x^m \in [-\infty, +\infty]$

$$\hat{y} = w_1 * x^1 + w_2 * x^2 + \dots + w_m * x^m + b$$

Diagram illustrating the mapping of feature ranges to the linear model coefficients:

- $[0, 100]$ maps to w_1
- $[0, 1]$ maps to w_2
- $[-\infty, +\infty]$ maps to w_m

$X \in \mathbb{R}^{n \times m}$

- Z-score normalization

- $X \leftarrow \frac{X - \mu}{\sigma}$

$\mu \in \mathbb{R}^{1 \times m}$
(average sample)


--	--	--	--	--	--

$\sigma \in \mathbb{R}^{1 \times m}$


--	--	--	--	--	--

Z-Score


Normalization




Raw data



Zero Mean



One Deviation



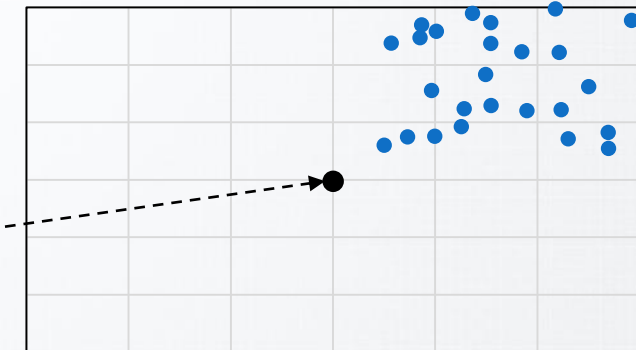
Z-Score

Z-Score


Normalization

Raw Data


Origin
($x^1 = 0, x^2 = 0$)




Zero Mean



One Deviation




Z-Score



Bibliography


Bibliography

- Dive into Deep Learning
 - Chapter 3
 - 3.1 Linear Regression
 - 3.4 Linear Regression Implementation from Scratch



Bibliography

- The Hundred-page Machine Learning Book
 - Chapter 3 – Fundamental Algorithms
 - 3.1 Linear Regression



Bibliography

- Deep Learning
 - Chapter 4 – Numerical Computation
 - 4.3 Gradient-Based Optimization

