

Recurrent Neural Networks

Prof. Artur Jordão

Definition

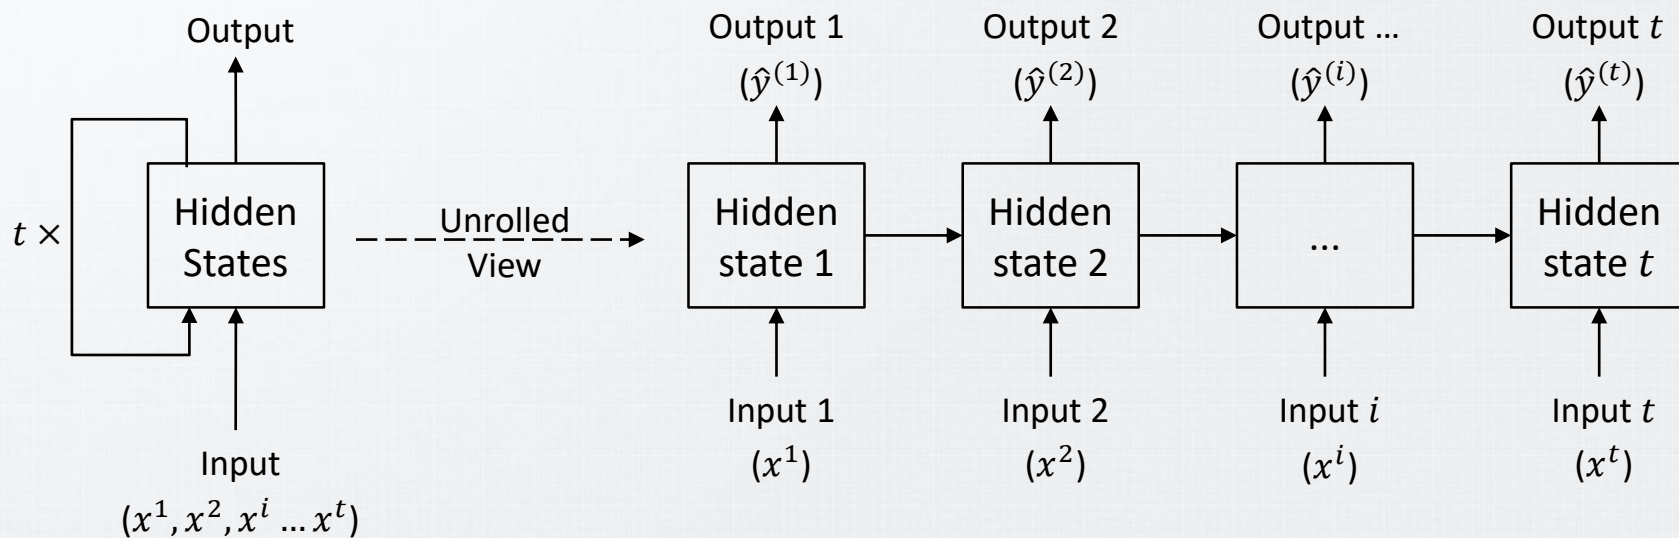
Recurrent Neural Networks

- Recurrent neural networks (RNNs) are models that capture the dynamics of **sequences** via recurrent connections
- Recurrent neural networks differ from feedforward networks (e.g., MLPs) by allowing cycles in their computation graphs
- RNNs are able to handle sequential and temporal data
 - They remember earlier parts of the sequence to interpret or contextualize later elements when making predictions

Architectural Design

Recurrent Neural Networks



- Recurrent neural networks unroll across time steps (or sequence steps), applying the **same underlying parameters** at each step
 - **The weights are shared across all time steps**



Hidden Layers vs. Hidden States

Recurrent Neural Networks

- Hidden layers
 - Layers that are hidden from view on the path from input to output
 - Layers where the training data doesn't reveal the desired output
- Hidden states
 - Inputs for a given step that can only be computed by looking at data from previous time steps
 - The hidden state at any time step: $h_t = f(x^t, h_{t-1})$



Hidden Layers vs. Hidden States

Recurrent Neural Networks

- Recurrent neural networks are neural networks with **hidden states**
 - Inputs from earlier time steps influence the RNN's response to the current input
 - Store all the data it has observed (memory)
- Causal structure
 - The state at time t captures information from the past: $x^{(1)}, \dots x^{(t-1)}$ as well as the current input $x^{(t)}$

Connections


Recurrent Neural Networks

- The RNN has the following connections
 - Input to hidden
 - Hidden to hidden
 - Hidden to output
- Input to hidden
 - Connections parameterized by a weight matrix U
- Hidden to hidden
 - Recurrent connections parameterized by a weight matrix W
- Hidden to output
 - Connections parameterized by a weight matrix V

Forward Propagation

Recurrent Neural Networks

- Forward propagation starts with specifying the initial state $h^{(0)}$
- For each time step from $t = 1$ to $t = \tau$
 - $a^{(t)} = b + Wh^{(t-1)} + Ux^{(t)}$
 - $h^{(t)} = \tanh(a^{(t)})$
 - $o^{(t)} = c + Vh^{(t)}$
 - $\hat{y}^{(t)} = \text{softmax}(o^{(t)})$
- The matrices **U, W, V** are shared across all time steps
 - Vanishing and exploding gradient problem



Loss Function


Recurrent Neural Networks

- The total loss for a given sequence of t ($x^{(1)}, \dots x^{(t)}$) values paired with a sequence of y values would then be just the sum of the losses over all the time steps
 - Assuming a recurrent network that maps an input sequence to an output sequence of the same length
- $$L(\{x^{(1)}, \dots x^{(t)}\}, \{y^{(1)}, \dots y^{(t)}\}) = \sum_{i=0}^t L^{(i)}$$
- Computing the gradient of this loss function w.r.t the parameters is computationally expensive
 - The runtime (and memory cost) is $O(t)$
 - Parallelization cannot reduce this cost because the forward propagation graph is **inherently sequential**: each time step must be computed after the previous one

Architectural Designs


Recurrent Neural Networks

- Recurrent networks that produce an output at each time step
 - Many-to-many
- Recurrent networks that read an entire sequence and then produce a single output
 - Many-to-one



Architectural Design

Recurrent Neural Networks



Variants of Recurrent Networks

Introduction

Variants of RNNs

- RNNs are known to forget information, preventing the modeling of **long-range relationships**
- The problems of learning long-term dependencies
 - Vanishing and exploding gradients
- Memory cell
 - It retains information over time
- Gated RNNs
 - Long Short-Term Memory (LSTM)
 - Gated Recurrent Units (GRU)

Long Short-Term Memory (LSTM)

Variants of RNNs

- Gating units
 - Vectors that control the flow of information in the LSTM via element-wise multiplication of the corresponding information vector
- The values for the gating units are always in the range $[0,1]$ and are obtained as the outputs of a sigmoid function applied to the current input and the previous hidden state

Long Short-Term Memory (LSTM)

Variants of RNNs

- Input Gate (I^t)
 - Determines how much of the input node's value should be added to the current memory cell's internal state
- Forget Gate (F^t)
 - Determines whether each element of the memory cell is remembered (copied to the next time step) or forgotten (reset to zero)
- Output Gate (O^t)
 - Determines whether the memory cell should influence the output at the current time step

Long Short-Term Memory (LSTM)

Variants of RNNs

- The update is (the bias is omitted for simplicity):

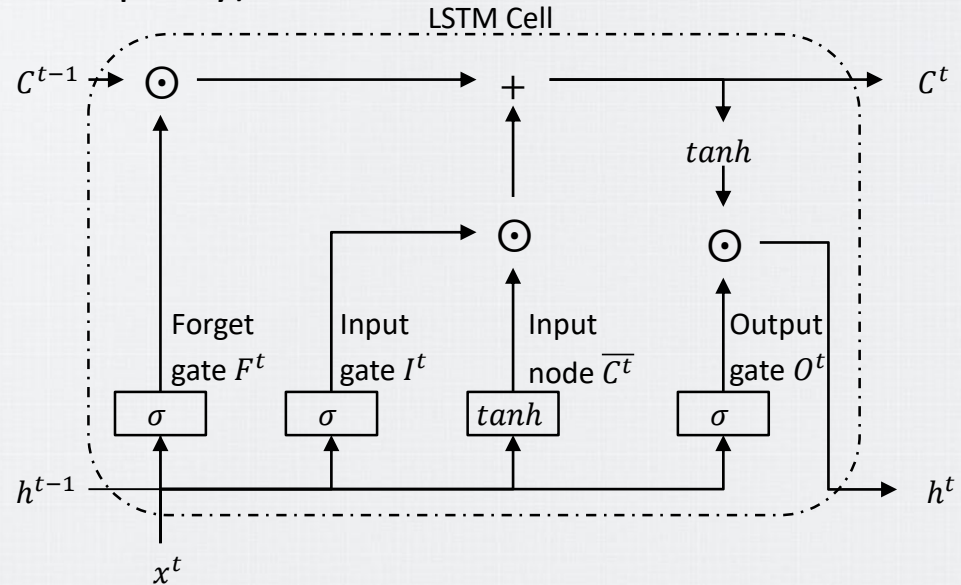
- $I^t = \sigma(W_I x^t + W_{hI} h^{(t-1)})$
- $F^t = \sigma(W_F x^t + W_{hF} h^{(t-1)})$
- $O^t = \sigma(W_O x^t + W_{hO} h^{(t-1)})$

- Input node

- $\bar{C}^t = \tanh(W_C x^t + W_{hC} h^{(t-1)})$

- Memory Cell Internal State

- $C^t = F^t \odot C^{(t-1)} + I^t \odot \bar{C}^t$



Long Short-Term Memory (LSTM)

Variants of RNNs

- The update is (the bias is omitted for simplicity):

- $I^t = \sigma(W_I x^t + W_{hI} h^{(t-1)})$
- $F^t = \sigma(W_F x^t + W_{hF} h^{(t-1)})$
- $O^t = \sigma(W_O x^t + W_{hO} h^{(t-1)})$

- Input node

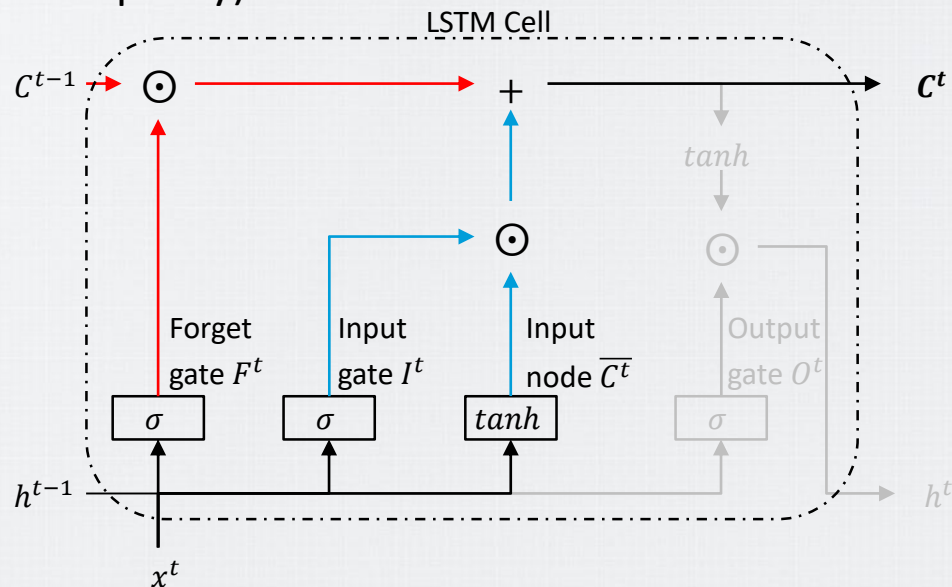
- $\bar{C}^t = \tanh(W_c x^t + W_{hc} h^{(t-1)})$

- Memory Cell Internal State

- $$C^t = \underbrace{F^t \odot C^{(t-1)}}_{\text{addresses how much of the old cell internal state } C^{(t-1)} \text{ we retain}} + \underbrace{I^t \odot \bar{C}^t}_{\text{governs how much we take new data into account via } \bar{C}^t}$$

addresses how much of the old cell internal state $C^{(t-1)}$ we retain

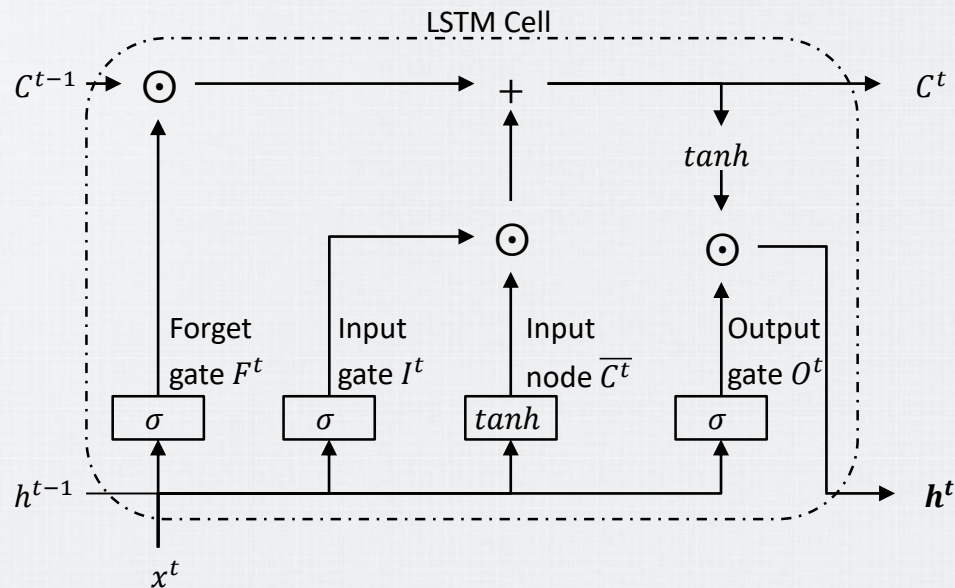
governs how much we take new data into account via \bar{C}^t



Long Short-Term Memory (LSTM)

Variants of RNNs


- Hidden State
 - $h^t = O^t \odot \tanh(C^t)$



Bibliography


Bibliography

- Deep Learning
 - Chapter 10
 - 10.2 Recurrent Neural Network



Bibliography

- Dive into Deep Learning
 - Chapter 9
 - 9.4.2 Recurrent Neural Networks with Hidden States



Bibliography

- The Hundred-page Machine Learning Book
 - Chapter 6
 - 6.2.2 Recurrent Neural Network

